

Kubernetes Configurations and Advance Pod Concepts

- starts at 9:05 pm

AGENDA

1. Kubectl apply summary
2. Labels and Annotations
3. Multi-container Pods and use Cases
4. Init Containers
5. Replication Controller and ReplicaSets
6. Deployments Introduction
7. DaemonSets in Kubernetes
8. Configmaps
9. Secrets
10. Inject Data into Applications
11. Resource Settings for pods in Kubernetes
12. Service Accounts
13. Pod Scheduling and Placement
14. Taints and Tolerations
15. Node Affinity/ Anti Affinity
16. Node Selectors and Labels

Kubectl Apply summary. → kubectl apply -f pod.yaml

① Kubectl client.



② API server.



③ etcd (podspec)



④ controller manager. → (API server create pod (pending))



⑤ Scheduler → updates podspec with the node it has decided.
(etcd)✓



⑥ Kubelet

The **Kubelet** on the assigned node continuously watches the API server for Pods scheduled to run on its node.

↓ Pulls pod spec from API server.
→ starts the pod.

⑦ CRI (pulls the image)
→ Run the container.

⑧ Kube proxy. (Networking)

Summary: Flow of Events

1. **kubectl apply** → Reads `kubeconfig` to authenticate with the API server and sends the Pod spec.
2. **API Server** validates the request, writes the desired state to **etcd**, and notifies other components.
3. **Controller Manager** creates the Pod in the API server (Pending state).
4. **Scheduler** updates the Pod spec in **etcd** with the node it has selected for the pod, and informs the API server.
5. **Kubelet** on the assigned node fetches the Pod spec, prepares the environment, and starts the containers.
6. **Container Runtime** runs the containers, and **Kube-Proxy** ensures networking is properly configured.

keeps monitoring the pod.

Labels and Annotations

① Labels

[pod]

env: prod //

→ Use Case

① Selection.

→ Selectors



(deployments,
Replicaset)

② Filtering.

kubectl get pods -l env:prod.

③ grouping.

labels:

key

app: my-app

value

environment: production

tier: backend

Annotations.

key : value.

eg: → giving description.

→ Automation.

Jenkins, argoCD.

annotations:

build: "v1.2.3"

description: "Pod running a production application"

maintainer: "team@mycompany.com"

pod.yaml

apiVersion: v1

kind: Pod

metadata:

name: demo-pod

labels:

app: demo-app

environment: dev

team: devops

annotations:

description: "This is a demo pod for labels and annotations"

build: "v1.0"

spec:

containers:

- name: nginx-container

image: nginx

****Command to show labels of the pod****

kubectl get pods --show-labels

****You can use `kubectl label` to add labels to your resources.****

kubectl label pod pod-name app=nginx

****Filtering pods based on labels****

kubectl get pods -l environment=production

****Label nodes****

kubectl label node node-name region=us-west


****Get nodes based on labels****

kubectl get nodes -l role=frontend

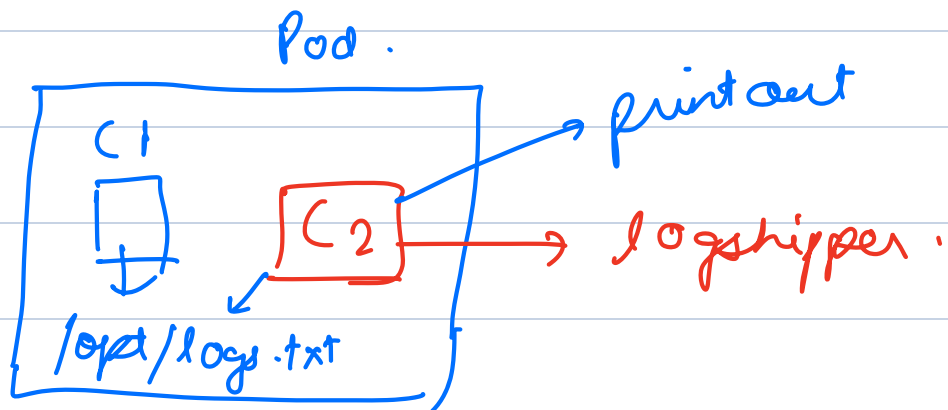
****Adding Annotations to a Pod****

kubectl annotate pod pod-name git-commit-hash=abc1234

Multi container Pods



① Sidecar Container.



apiVersion: v1

kind: Pod

metadata:

name: myapp

labels:

app: myapp

spec:

containers:

- name: myapp

image: alpine:latest

command: ['sh', '-c', 'while true; do echo "logging \$(date)" >> /opt/logs.txt; sleep 1; done']

volumeMounts:

- name: data

mountPath: /opt

- name: logshipper

image: alpine:latest

command: ['sh', '-c', 'tail -F /opt/logs.txt']

volumeMounts:

- name: data

mountPath: /opt

volumes:

- name: data

emptyDir: {}

kubectrl logs -f myapp -c logshipper

kubectrl exec -it myapp -c myapp -- sh

→ init Containers.

① Initialising Tasks →
→ Downloading

② Preflight Checks.
→ verifying DB connection
→ checking for files.

apiVersion: v1

kind: Pod

metadata:

name: init-demo

spec:

initContainers:

- name: init-container

image: alpine:latest

command: ['sh', '-c', 'echo "Initializing the app..." > /data/init.txt; sleep 5']

volumeMounts:

- name: shared-data

mountPath: /data

containers:

- name: main-container

image: alpine:latest

command: ['sh', '-c', 'while true; do echo "Main container is running"; sleep 1; done']

volumeMounts:

- name: shared-data

mountPath: /data

volumes:

- name: shared-data

emptyDir: {}

Break → 10:20 pm

→ Replication Controller & Replicaset

apiVersion: v1

kind: ReplicationController

metadata:


name: my-replication-controller

spec:

replicas: 3

selector:

app: my-app



template:

metadata:

labels:

app: my-app

spec:

containers:

- name: my-container


image: nginx


RC

image

container

pod





Feature	ReplicationController	ReplicaSet
Label Selector	Supports equality-based selectors (e.g., <code>key=value</code>).	Supports both equality-based and set-based selectors (e.g., <code>key in (value1, value2)</code>).
Flexibility	Less flexible in managing pods.	More flexible due to advanced label selectors.
Integration with Deployments	Not used by Deployments.	Used as the underlying mechanism for Deployments.
Purpose	Basic functionality for managing pod replicas.	Enhanced functionality for modern Kubernetes workflows.
Deprecation	Considered legacy and no longer recommended for new setups.	Actively supported and preferred.

selector:

} //

app: nginx

selector:

matchExpressions:

} → set based selector.

- key: app

app = [nginx, apache]

operator: In

values:

- nginx

- apache

Pod 1 app = nginx

Pod 2 app = apache.

1. ****Ensures Desired State****: It ensures that a specific number of pod replicas are always running.

2. ****Supports Set-based Selectors****: It can filter pods using labels with set-based selectors.

3. ****Scaling****: You can scale the number of replicas up or down.

4. ****Self-healing****: If pods fail, it ensures new ones are created to replace them. Increasing or decreasing the number of replicas).

→ Demo

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: my-replicaset

spec:

replicas: 3

selector:

matchExpressions:

- key: app

operator: In

values:

- my-app

template:

metadata:

labels:

app: my-app

spec:

containers:

- name: nginx-container

image: nginx

ports:

- containerPort: 80

kubectl apply -f replicaset.yaml

kubectl get rs

kubectl get pods | grep replic

kubectl delete pod (any 1 pod)

****Scale Out ReplicaSet****

kubectl scale rs my-replicaset --replicas=5

****Scale in ReplicaSet****

kubectl scale rs my-replicaset --replicas=1

Limitations

→ No rolling updates

update image in replicaset → X

→ No rollbacks.

→ Deployments

→ Rolling update


→ Rollbacks.


Deployments provide a higher abstraction layer, enabling features like rolling updates, rollbacks, and version management while relying on ReplicaSets under the hood.

A **Deployment** is an abstraction that manages ReplicaSets and ensures the desired state of your applications. It is the most common way to run applications in Kubernetes and adds rolling update functionality and rollback features.

How Deployment Works

1. **Pod Creation**: Kubernetes will create a new ReplicaSet under the hood when you create a Deployment. The ReplicaSet will manage the pods according to the `template`.
2. **Rolling Updates**: When the Deployment is updated (e.g., a new container image is provided), Kubernetes performs a rolling update by gradually replacing old pods with new ones.
3. **Rollbacks**: If something goes wrong with the update, you can rollback to the previous version of the Deployment.
4. **Scaling**: Similar to ReplicaSets, you can scale the number of pods by adjusting the `replicas` field in the Deployment.





apiVersion: apps/v1

kind: Deployment

metadata:

name: my-deployment

spec:

replicas: 3

selector:

matchLabels:

app: my-app

template:

metadata:

labels:

app: my-app

spec:

containers:

- name: nginx-container

image: nginx:1.23

ports:

- containerPort: 80

kubectl apply -f deployment.yaml

****Update the Deployment**:**

kubectl set image deployment/my-deployment nginx-container=nginx:1.24

kubectl get pods -w | grep my-deployment


****Rollback the Deployment**:**

kubectl rollout undo deployment/my-deployment

****Rollback history**** - View the history of rollouts for `my-deployment`

kubectl rollout history deployment/my-deployment

→ Daemonsets



logging agents
monitoring
networking components

nodeSelector:

role = storage.

role: storage

DS → YAML

apiVersion: apps/v1

kind: DaemonSet

metadata:

name: nginx-daemonset

spec:

selector:

matchLabels:

name: nginx

template:

metadata:

labels:

name: nginx

spec:

containers:

- name: nginx

image: nginx:1.14.2

ports:

- containerPort: 80

with nodeselector.

apiVersion: apps/v1

kind: DaemonSet

metadata:

name: custom-nginx-daemonset

spec:

selector:

matchLabels:

name: nginx

template:

metadata:

labels:

name: nginx

spec:

nodeSelector:

kubernetes.io/hostname: new-cluster1-worker2

containers:

- name: nginx

image: nginx:1.14.2

ports:

- containerPort: 80