

Kubernetes Observability and Pod Design

Continued

Starts at 9:05 pm

Agenda

1. Node Selector and Labels
2. Affinity and Anti Affinity
3. Probes in kubernetes
4. Metrics Server
5. Deployment Strategies

→ Node selectors and Labels

Nodes → labels

nodeselector → pod spec.

```
kubectl label nodes new-cluster1-worker disktype=ssd region=us-west
```

```
kubectl label nodes new-cluster1-worker2 disktype=hdd region=us-east
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

name: multi-label-selector-pod

spec:

nodeSelector:

disktype: ssd

region: us-west

containers:

- name: example-container

image: nginx

command: ["sh", "-c", "echo Hello from us-west SSD node!"]

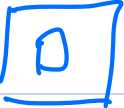
****Removing Labels from the nodes****

kubectl label nodes new-cluster1-worker2 disktype- region-

****Removing Labels from the nodes****


kubectl label nodes new-cluster1-worker disktype- region-

→ Affinity and Anti Affinity

us-west ✓


India .


- ① Node Affinity → Matching labels.
- ② Pod Affinity
- ③ Pod Anti-affinity.



****Node Affinity****: Rules to ensure pods are scheduled on nodes with specific labels.

- ****Pod Affinity****: Rules to co-locate pods near other pods based on their labels.

- ****Pod Anti-Affinity****: Rules to prevent pods from being scheduled on the same nodes as other pods with specific labels.

Node Affinity.

1. ****RequiredDuringSchedulingIgnoredDuringExecution****:

- The pod ****can only be scheduled**** on nodes that ****match**** the affinity rules specified.

- If no node matches the rules, the pod will ****not be scheduled****.

- This is a ****hard requirement**** for scheduling.

1. ****PreferredDuringSchedulingIgnoredDuringExecution****:

- The pod ****prefers**** to be scheduled on nodes that ****match**** the affinity rules specified, but it is not a hard requirement.

- If no nodes match, the pod can still be scheduled on any available node.

- This is a ****soft preference**** that Kubernetes will try to honor but will not block scheduling if no matching node is found.

apiVersion: v1

kind: Pod

metadata:

name: my-pod

spec:

affinity:

nodeAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

nodeSelectorTerms:

- matchExpressions:

- key: kubernetes.io/hostname

operator: In

values:

- new-cluster1-worker

- new-cluster1-worker2

preferredDuringSchedulingIgnoredDuringExecution:

- weight: 1

preference:

matchExpressions:

- key: environment

operator: In

values:

- production

containers:

- name: my-container

image: nginx

Pod affinity.

Pod Affinity is a concept in Kubernetes that allows you to schedule pods onto nodes based on the **labels** of other pods running on the same or different nodes.

This is useful in scenarios where you want to keep related pods together, for example:

- Pods in the same application stack should run on the same node to reduce latency.

- Running high-throughput services together on the same node to improve network performance.

apiVersion: v1

kind: Pod

metadata:

name: my-pod

```
spec:

  affinity:

    podAffinity:

      requiredDuringSchedulingIgnoredDuringExecution:

        - labelSelector:

            matchLabels:

              app: my-app

            topologyKey: kubernetes.io/hostname

    containers:

      - name: my-container

        image: nginx
```

Pod Anti Affinity

Pod anti-affinity ensures that pods are ****spread out**** and not scheduled on the same nodes as other pods with specific labels. Useful for:

```
apiVersion: v1

kind: Pod

metadata:

  name: pod-anti-affinity-pod

  labels:

    app: my-app
```

spec:

affinity:

podAntiAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

- labelSelector:

matchLabels:

app: my-app

topologyKey: "kubernetes.io/hostname"

containers:

- name: nginx

image: nginx

Combined use

apiVersion: v1

kind: Pod

metadata:

name: combined-affinity-pod

spec:

affinity:

nodeAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

nodeSelectorTerms:

- matchExpressions:

- key: disk-type

operator: In

values:

- ssd

podAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

labelSelector:

matchLabels:

app: backend

topologyKey: "kubernetes.io/hostname"

podAntiAffinity:

preferredDuringSchedulingIgnoredDuringExecution:

- weight: 1

podAffinityTerm:

labelSelector:

matchLabels:


app: frontend

topologyKey: "kubernetes.io/hostname"

containers:



- name: nginx

image: nginx




→ Probes

1. **Liveness Probe** – Checks if the container is still running. If the probe fails, Kubernetes restarts the container.
2. **Readiness Probe** – Checks if the container is ready to serve traffic. If the probe fails, the Pod is removed from the Service endpoints.
3. **Startup Probe** – Used for slow-starting applications. It prevents Kubernetes from marking a container as failed before it has fully started.



Deployment



→ endpoint (backend = deployment)
↓
Service.

Break → 10:21 pm

→ Liveness Probe

http://< pod-ip > : 80 / health

↓
404

http://< pod-ip > : 80

↓

Liveness: http-get http://:80/ delay=5s timeout=1s period=10s #success=1 #failure=3

| Parameter | Explanation |
|-------------------------|---|
| http-get http://:80/ | The probe sends an HTTP GET request to port 80 of the container to check if it's alive. |
| delay=5s | The first probe starts 5 seconds after the container begins running (initial delay). |
| timeout=1s | The probe waits 1 second for a response before considering it a failure. |
| period=10s | The probe runs every 10 seconds to check the container's health. |
| #success =1 | (Not typically seen in the liveness probe) If used, it would mean the container is considered healthy after one successful probe. |
| #failure =3 | If the probe fails 3 times consecutively, Kubernetes restarts the container. |

livenessProbe:

httpGet:

path: /

port: 80

initialDelaySeconds: 5

timeoutSeconds: 2

periodSeconds: 10

failureThreshold: 5 # Custom value, default is 3

successThreshold: 2 # Custom value, default is 1

Use Cases of Liveness Probe

- **Crash Recovery**: If a container crashes due to an application error or other issue, the liveness probe can detect this and restart the container automatically.

- **Deadlocks and Stuck Processes**: If an application inside a container becomes unresponsive due to a deadlock or gets stuck, the liveness probe can restart the container to recover from this state.

- **Ensuring High Availability**: By continuously monitoring the health of containers, liveness probes help maintain the high availability and reliability of applications.

Use Cases of Readiness Probe

- **Rolling Updates**: During rolling updates, readiness probes help ensure that only the updated and ready instances of a container serve traffic. If a container fails the readiness check, it won't receive any traffic until it passes the check.
- **Graceful Shutdowns**: Readiness probes can be used to remove a container from service endpoints before it is terminated, ensuring that no new traffic is sent to a container that is shutting down.

Readiness Probe

apiVersion: v1

kind: Pod

metadata:

name: liveness-readiness-pod

spec:

containers:

- name: my-app

image: nginx

ports:

- containerPort: 80

livenessProbe:

httpGet:

path: /health

port: 80

initialDelaySeconds: 5

periodSeconds: 10

readinessProbe:

httpGet:

path: /health

port: 80

initialDelaySeconds: 3

periodSeconds: 5

Startup Probes.

Kubernetes waits for the startup probe to succeed before performing any liveness or readiness probes.

Startup Probe Fields:

- **httpGet**: Performs an HTTP GET request to a specified path and port.
- **tcpSocket**: Checks whether the container is listening on a specific port.
- **exec**: Runs a command inside the container to check if the application is ready.

apiVersion: v1

kind: Pod

metadata:

name: my-pod

spec:

containers:

- name: mysql

image: mysql:5.7

env:

- name: MYSQL_ROOT_PASSWORD

value: rootpassword

startupProbe:

exec:

command:

- "mysql"

- "-uroot"

- "-prootpassword"

- "-e"

- "SELECT 1;"

initialDelaySeconds: 60 # Allow MySQL enough time to start

periodSeconds: 5

failureThreshold: 10

timeoutSeconds: 3

ports:

- containerPort: 3306

Use Cases

- ****Applications with Long Startup Times****: Some applications may take a long time to initialize. Startup probes ensure that these applications are not killed or marked as unready during their initialization phase.
- ****Preventing Premature Failures****: By delaying liveness and readiness checks until the application has started, startup probes prevent premature failures and restarts.
- ****Ensuring Smooth Startups****: Startup probes help in managing the startup sequence smoothly, especially for applications with complex initialization processes.

Final Summary: Liveness vs. Readiness vs. Startup Probes

| Probe Type | Purpose | Effect if Fails | Restarts Container? | Removes Pod from Service? | When to Use? |
|------------------------|---|--|---------------------|---------------------------|---|
| Readiness Probe | Checks if the pod is ready to serve traffic | No traffic sent to pod | ❌ No | ✅ Yes | When a pod needs time to become ready (e.g., waiting for DB connection) |
| Liveness Probe | Checks if the container is still running properly | Restarts the container | ✅ Yes | ❌ No | When a container may get stuck or become unresponsive |
| Startup Probe | Checks if the container has finished its startup sequence | Kills and restarts the container if it doesn't start in time | ✅ Yes | ❌ No | When an application has a long startup time before becoming healthy |

Metrics Server.

CPU RAM nodes and pods.

wget https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml

spec:

containers:

- name: metrics-server

args:

- --cert-dir=/tmp

- --secure-port=4443

- --kubelet-insecure-tls # Add this line to bypass certificate errors

kubectl top pod --all-namespaces

kubectl top node

→ Deployment strategies

① Rolling update.

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

spec:

replicas: 3

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx:1.21

ports:

- containerPort: 80

strategy:

type: RollingUpdate

rollingUpdate:



maxSurge: 2 # Allows two additional pod

maxUnavailable: 2 # Allows two pod to be unavailable





max surge $\rightarrow 2$





unavailable $\rightarrow 2$




- **maxSurge**: This controls how many additional pods (above the desired replica count) can be created during an update. It allows extra pods to be spun up temporarily during the update.



- **maxUnavailable**: This controls how many pods (below the desired replica count) can be unavailable during an update. This ensures there is no sudden loss of available pods.

Blue - green Deployment .





- 2 CMS. → green.
- Blue.
- 2 deployments → green.
- Blue.
- 1 service.



→ 8080

curl machine:8080

****Create configmaps****

```
kubectl create configmap nginx-blue --from-literal=index.html='<h1>Blue</h1>'
```

```
kubectl create configmap nginx-green --from-literal=index.html='<h1>Green</h1>'
```

****Blue Deployment****

apiVersion: apps/v1

kind: Deployment

metadata:

name:

name: my-app-blue

spec:

replicas: 3

selector:

matchLabels:

app: my-app

version: blue

template:

metadata:

labels:

app: my-app

version: blue

spec:

containers:

- name: nginx

image: nginx:1.23

volumeMounts:

- name: nginx-config

mountPath: /usr/share/nginx/html

volumes:

- name: nginx-config

configMap:

name: nginx-blue

kubectl expose deployment my-app-blue --port=80 --target-port=80 --name=my-app-service

kubectl port-forward svc/my-app-service 8080:80

****Create Green Deployment****

```
apiVersion: apps/v1

kind: Deployment

metadata:

  name: my-app-green

spec:

  replicas: 3

  selector:

    matchLabels:

      app: my-app

      version: green

  template:

    metadata:

      labels:

        app: my-app

        version: green

    spec:

      containers:

        - name: nginx

          image: nginx:1.24

          volumeMounts:

            - name: nginx-config

              mountPath: /usr/share/nginx/html
```

volumes:

- name: nginx-config

configMap:

name: nginx-green


****Patch the service to update the labels.****

```
kubectl patch service my-app-service -p '{"spec":{"selector":{"app":"my-app","version":"green"}}}'
```


****Again do port forwarding****

```
kubectl port-forward svc/my-app-service 8080:80
```

Canary Deployment.



→ v2



create 1 deploy. (stable)


3 replicas

create 1 deploy (canary)

1 replica.

create service.

→ app = my-app.



Scale stable = 1

Scale canary = 3

Region 1

Region 2

N1

N2

N3

N4

\square

D

\square

\square

