

# Kubernetes Jobs and Networking

starts at 9:05 pm

## Agenda

### 1. Jobs

#### 1. Completions and Parallelism

### 2. Cron jobs

### 3. Services

#### 1. Types of Services

### 4. Network Policies

## → Jobs

#### \*\*Backup Tasks\*\*

- Periodic backups of a database (e.g., MySQL, PostgreSQL)

- Exporting logs or data snapshots to cloud storage

#### - \*\*Batch Processing\*\*

- Generating reports (e.g., monthly sales reports)

- Running analytics jobs on collected data

#### - \*\*Cleanup Jobs\*\*

- Deleting old logs, cache, or expired records

- Cleaning up orphaned resources

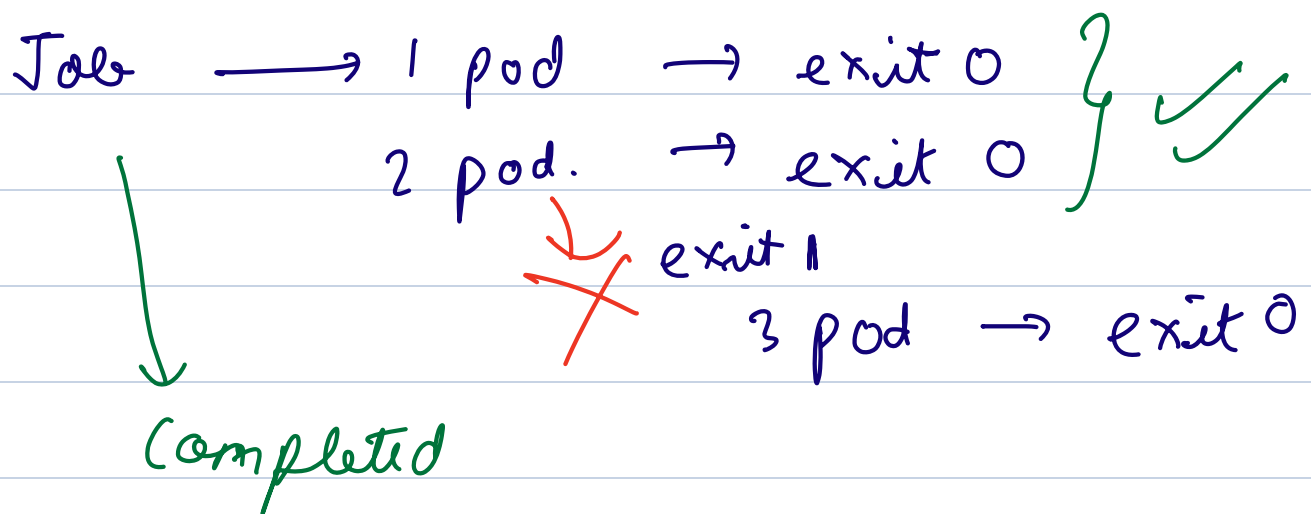
#### - \*\*Data Import/Export\*\*

- Fetching external data and importing it into a database

- Sending bulk emails or notifications
- **Security Scans & Compliance Checks**
- Running vulnerability scans on application containers
- Checking compliance rules on cloud resources

## Jobs

In Kubernetes, **jobs** and **cron jobs** are essential for running tasks that are not part of the main application but need to be executed periodically or as one-off tasks.



backoff limit → 4

Container fails 4 times

Job → Failed

## → Completion and Parallelism

how many pods need to run?  
how many times the job needs to  
complete successfully?

Completion →

Defines how many successful runs are required before the job is considered complete.

default → 1

Parallelism →

→ how many pods can run at a time  
default → 1

spec:

completions: 5

parallelism: 1

2 →

2 pods.

2 pods

1 pod

} Completed

spec:

completions: 10

parallelism: 5

Scenario	Completions	Parallelism	Use Case
One-time job	1	1	Database migration, cleanup tasks
Fixed number of runs (sequential)	N	1	Data processing in sequence
Fixed number of runs (parallel)	N	M	Large-scale computations

#### #### Key Parameters:

- `completions`: The desired number of successful pod completions.
- `parallelism`: The number of pods to run in parallel.
- `backoffLimit`: The number of retries before the job is considered failed.
- `activeDeadlineSeconds`: The time duration before the job is considered failed, even if it's still running.

#### ### Case 1 backofflimit

apiVersion: batch/v1

kind: Job

metadata:

name: demo-job

spec:

completions: 5 # Total 5 successful runs required

parallelism: 2 # Run 2 pods at a time

backoffLimit: 3 # Retry a failed pod 3 times before marking job as failed

activeDeadlineSeconds: 60 # Job must finish within 60 seconds

template:

spec:

containers:

- name: worker

image: busybox

command: ["sh", "-c", "echo Running task; sleep 10; exit 1"]

restartPolicy: Never

apiVersion: batch/v1

kind: Job

metadata:

name: demo-job

spec:

completions: 5 # Total 5 successful runs required

parallelism: 2 # Run 2 pods at a time

backoffLimit: 3 # Retry a failed pod 3 times before marking job as failed

activeDeadlineSeconds: 10 # Job must finish within 60 seconds

template:

spec:

containers:

- name: worker

image: busybox

command: ["sh", "-c", "echo Running task; sleep 10; exit 0"]

restartPolicy: Never

r

apiVersion: batch/v1

kind: Job

metadata:

name: demo-job

spec:

completions: 5 # Total 5 successful runs required

parallelism: 2 # Run 2 pods at a time

backoffLimit: 3 # Retry a failed pod 3 times before marking job as failed

activeDeadlineSeconds: 60 # Job must finish within 60 seconds

template:

spec:

containers

name: worker

image: busybox

command: ["sh", "-c", "echo Running task; sleep 10; exit 0"]

restartPolicy: Never

# Cron Job

apiVersion: batch/v1

kind: CronJob

metadata:

name: demo-cronjob

spec:

schedule: "\*/2 \* \* \* \*" # Runs every 2 minutes

jobTemplate:

spec:

completions: 5 # Each Job requires 5 successful pods

parallelism: 2 # Run 2 pods at a time

backoffLimit: 3 # Retry a failed pod 3 times

activeDeadlineSeconds: 60 # Each job must complete in 60 seconds

template:

spec:

containers:

- name: worker

image: busybox

command: ["sh", "-c", "echo Running task; sleep 10"]

restartPolicy: Never

Cronjob → schedule.



↓ 2 mins.

→ Job → 60s

```
kubectl patch cronjob demo-cronjob -p '{"spec": {"suspend": true}}'
```

Break → 10:20 pm

→ Services



a **Service** is an abstraction that provides **stable networking** for a group of Pods.

Office Building → K8s cluster

multiple employees → Pods.

↓

desk phone → IP



helpline number → Service.

Route ↓

current employee

↓ Route

Pods

Types of services.

① Cluster IP

→ Used for communication only within the cluster

metadata:

labels:

app: myapp

} → Pod.

selector:

app: myapp


} → Service.

→ Pod

→ CM

→ Service

→ nginx-pod.



```
kubectl port-forward resource LOCAL_PORT:TARGET_PORT
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-pod
```

```
  labels:
```

```
    app: nginx
```

```
spec:
```

```
  containers:
```

```
    - name: nginx
```

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: html

mountPath: /usr/share/nginx/html

volumes:

- name: html

configMap:

name: nginx-html

---

apiVersion: v1

kind: ConfigMap

metadata:

name: nginx-html

data:

index.html: |

<html><body><h1>Welcome to My Nginx Pod</h1></body></html>

---

apiVersion: v1

kind: Service

metadata:

name: nginx-service

spec:

selector:

app: nginx # Matches the pod label

ports:

- protocol: TCP

port: 80

targetPort: 80

type: ClusterIP # Internal service

Alternative way to expose → Nodeport.

nodeport range. 30000 - 32767

Receptionist → Nodeport.

http://Node\_IP:NodePort

ports:

- protocol: TCP

port: 80

targetPort: 80

nodePort: 30007

→ curl



creates nodeport on your node.



:80 of the service.



target port 80 → on the container.

Demo.

3 pods.      3 CMs.  
→ NP service.

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod-1

labels:

app: nginx

pod: "1"

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: html

mountPath: /usr/share/nginx/html

volumes:

- name: html

configMap:

name: nginx-html-1

---

apiVersion: v1

kind: ConfigMap

metadata:

name: nginx-html-1

data:

index.html: |

<html><body><h1>Welcome to Pod 1</h1></body></html>

---

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod-2

labels:

app: nginx

pod: "2"

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: html

mountPath: /usr/share/nginx/html

volumes:

- name: html

configMap:

name: nginx-html-2

---

apiVersion: v1

kind: ConfigMap

metadata:

name: nginx-html-2

data:

index.html: |

```
<html><body><h1>Welcome to Pod 2</h1></body></html>
```

---

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod-3

labels:

app: nginx

pod: "3"

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: html

mountPath: /usr/share/nginx/html

volumes:



- name: html

configMap:

name: nginx-html-3

---

apiVersion: v1

kind: ConfigMap

metadata:

name: nginx-html-3

data:

index.html: |

<html><body><h1>Welcome to Pod 3</h1></body></html>

---

apiVersion: v1

kind: Service

metadata:

name: nginx-service

spec:

selector:

app: nginx # Selects all pods with the app label 'nginx'

type: NodePort # Exposes the service externally on every node

ports:

- protocol: TCP


port: 80 # The service port

targetPort: 80 # The port on the pod where traffic is forwarded

nodePort: 30007 # The external port to access the service

→ Drawbacks .

- ① Limited Port Ranges.
- ② Requires access to node IPs.
- ③ Security .
- ④ Not Proper Load Balancing .



→ Load Balancer Service .

## \*\*Step-by-Step Flow of NLB Provisioning in EKS\*\*

1. \*\*You define a Kubernetes `Service` of type `LoadBalancer`\*\*

- You include the annotation `service.beta.kubernetes.io/aws-load-balancer-type: "nlb"`

2. \*\*Kubernetes Cloud Controller Manager (CCM) detects the LoadBalancer request\*\*

- The \*\*EKS CCM\*\* (built into the control plane) watches for services of type `LoadBalancer`.

- It **communicates with AWS APIs** to provision an NLB.

### 3. **AWS Automatically Creates and Configures an NLB**

- An **AWS Network Load Balancer** is created.
- A **static external IP (Elastic IP)** is assigned if needed.
- The NLB forwards traffic to the Kubernetes **worker nodes** where the service's pods are running.

### 4. **Traffic is sent to pods**

- If `target-type: "ip"` is specified, NLB directly forwards traffic to the pod IPs instead of NodePorts.

apiVersion: v1

kind: Service

metadata:

name: my-nlb-service

annotations:

service.beta.kubernetes.io/aws-load-balancer-type: "nlb"

service.beta.kubernetes.io/aws-load-balancer-scheme: "internet-facing"

service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: "ip"

spec:

type: LoadBalancer

ports:

- port: 80

targetPort: 80

protocol: TCP

selector: \_\_\_\_\_

app: my-app \_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_

\_\_\_\_\_