

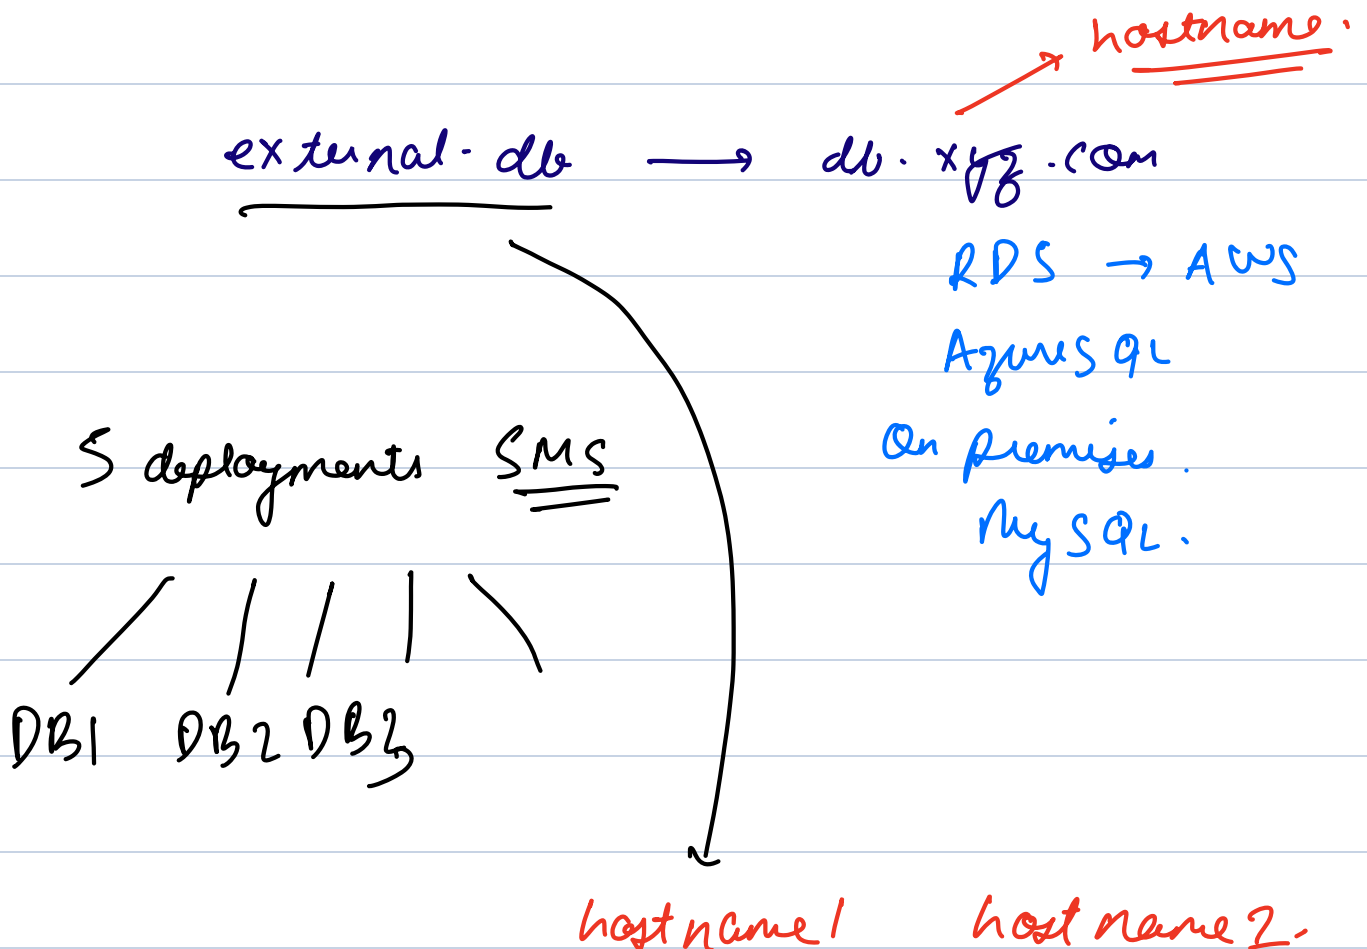
# Kubernetes Networking Continued

starts at 9:05 pm

## AGENDA

1. External Name Service
2. Headless Service
3. Ingress
4. Network Policies

### → External Name Service



apiVersion: v1

kind: Service

metadata:

name: external-db

spec:

type: ExternalName

externalName: example.com # Redirects traffic to example.com

kubectl run test-pod --image=busybox --restart=Never -it --rm -- /bin/sh

nslookup external-db

→ Headless service

No Cluster IP      No load balancing.

returns pod IPs.

→ stateful applications (DBs)

The **absence of an actual ClusterIP** tells Kubernetes **not to perform load balancing**

apiVersion: v1

kind: Service

metadata:

name: nginx-headless

spec:

selector:

app: nginx

nslookup nginx-headless.

clusterIP: None # This makes it a Headless Service

ports:

- protocol: TCP


port: 80

targetPort: 80

Best Practice for using services.

- ① Use selectors appropriately.
- ② Resource Management.
- ③ Security.

# Ingress.



→ Rules.

① hostname based

② Path - Based.

app.example.com → Service 1

api.example.com → service 2

/app1  
↓  
Service 1



/app 2  
↓  
Service 2.

amazon.com/spots

↓  
S1

amazon.com/Books.

↓  
S2



→ Set of Rules.



Ingress Controller → implements the ingress rules

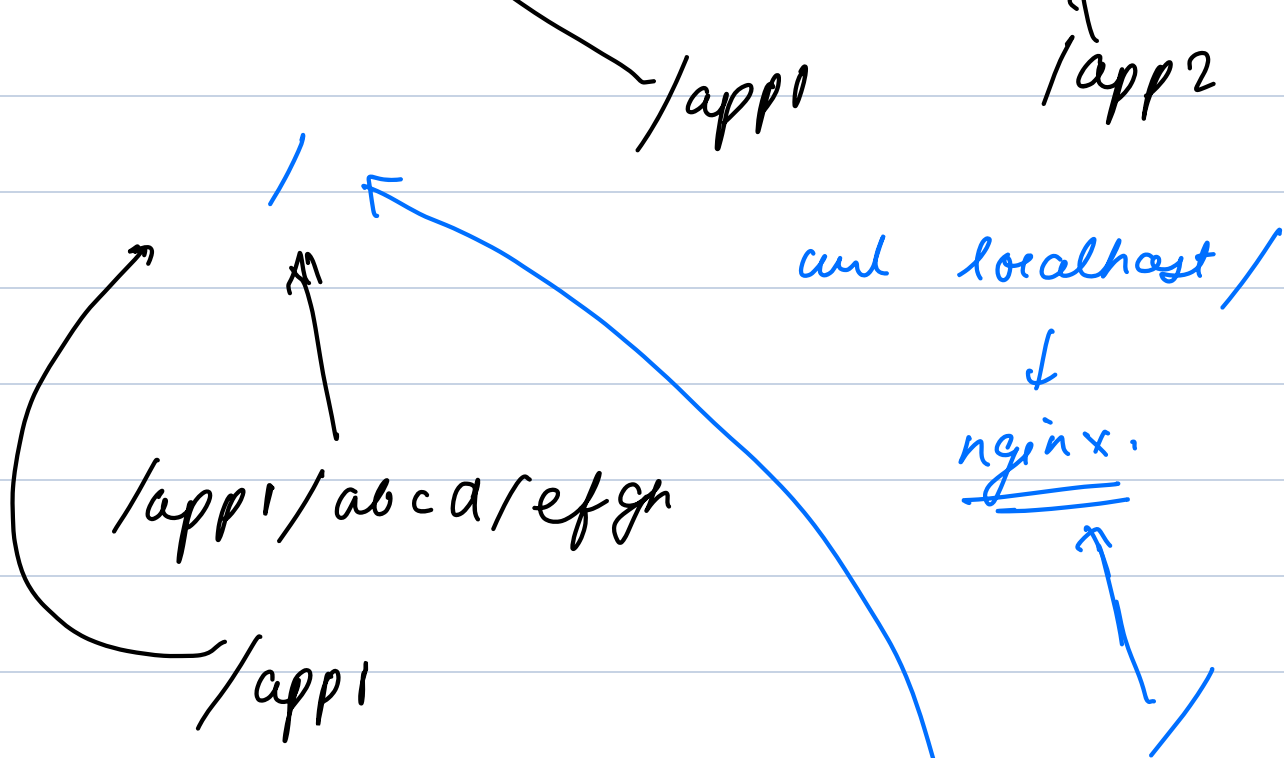
Think of **Ingress** as a blueprint for traffic routing and the **Ingress Controller** as the worker that follows that blueprint to direct traffic properly.

→ Demo.

- ① Create a new kind Cluster
- ② install nginx ingress controller
- ③ Create deployments and svc.
- ④ Create ingress Resource
- ⑤ Test.

local host: 80 → control-plane: 80





`/app2/abcd/efgh`

→ `my-app.local /app1`  
(Ingress)

`curl localhost/`  
↓  
`nginx.`

`/etc/hosts 2`

`my-app.local → 127.0.0.1`

my-app.local /app1

/app2  
↓

app2-service.  
↓

app1-service  
↓

app1-pod 127.0.0.1/  
↪

app2-pod.

1. `curl my-app.local/app1` sends a request.

→ EC2

2. `/etc/hosts` maps `my-app.local` to the IP of the `Kind node`.

→ Control Plane  
node: 80

3. The request reaches the `Ingress Controller` inside the cluster.

4. The Ingress Controller matches the `host` (`my-app.local`) and `path` (`/app1`) in the rule.

5. It routes the request to the `app1-service`.

6. `App1` processes the request and responds.

7. The response returns to `curl`.

↓  
VM: 80 →  
Node: 80

→ 10:35 Break.

→ curl localhost/app1 ?  
↓ )

curl localhost/

Ingress.  
↓

http://my-app.local/app1 → app1-servi  
↓

http://my-app.local/

kubectl apply -f <https://raw.githubusercontent.com/kubernetes/ingress-nginx/main/deploy/static/provider/kind/>

deploy.yaml

→ Install ingress Controller

kubectl get pods -n ingress-nginx

create deployment, CM, Service

apiVersion: v1

kind: ConfigMap

metadata:

name: app1-config

data:

index.html: |



```
<html>
```

```
<head><title>App 1</title></head>
```

```
<body style="background-color: lightblue;">
```

```
<h1>Welcome to App 1!</h1>
```

```
</body>
```

```
</html>
```

---

apiVersion: v1

kind: ConfigMap

metadata:

name: app2-config

data:

index.html: |

```
<html>
```

```
<head><title>App 2</title></head>
```

```
<body style="background-color: lightgreen;">
```

```
<h1>Welcome to App 2!</h1>
```

```
</body>
```

```
</html>
```

---

apiVersion: apps/v1

kind: Deployment

metadata:

name: app1

spec:

replicas: 1

selector:

matchLabels:

app: app1

template:

metadata:

labels:

app: app1

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: config-volume

mountPath: /usr/share/nginx/html

volumes:

- name: config-volume

configMap:

name: app1-config

---

apiVersion: v1

kind: Service

metadata:

name: app1-service

spec:

selector:

app: app1

ports:

- protocol: TCP

port: 80

targetPort: 80

---

apiVersion: apps/v1

kind: Deployment

metadata:

name: app2

spec:

replicas: 1

selector:

matchLabels:

app: app2

template:

metadata:

labels:

app: app2

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

volumeMounts:

- name: config-volume

mountPath: /usr/share/nginx/html

volumes:

- name: config-volume

configMap:

name: app2-config

---

apiVersion: v1

kind: Service

metadata:

name: app2-service

spec:

selector:

app: app2

ports:

- protocol: TCP

port: 80

targetPort: 80

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: my-ingress

annotations:

nginx.ingress.kubernetes.io/rewrite-target: /

spec:

ingressClassName: nginx

rules:

- host: my-app.local

http:

paths:

Create ingress Resource

- path: /app1

pathType: Prefix

backend:

service:

name: app1-service

port:

number: 80

- path: /app2

pathType: Prefix

backend:

service:

name: app2-service

port:

number: 80

sudo vi /etc/hosts


127.0.0.1 my-app.local

curl http://my-app.local/app1

curl http://my-app.local/app2

# Network Policies

Network Policies in Kubernetes define **rules for controlling** how Pods communicate with each other and external resources.



## Key Concepts of Network Policies

1. **Pod Selector**: A label selector to select which pods the policy applies to.

app = nginx

2. **Ingress and Egress Rules**:

- **Ingress**: Defines the allowed incoming traffic to a pod.

- **Egress**: Defines the allowed outgoing traffic from a pod.

3. **Namespace Selector**: Used to specify namespaces from which traffic is allowed.

4. **IP Blocks**: Allows you to specify traffic from specific IP ranges.

5. **Ports**: You can restrict traffic based on ports.

Ingress → Incoming traffic pod / svc  
Egress → Outgoing traffic

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: allow-frontend-to-backend

namespace: default

spec:

podSelector:

matchLabels:

app: backend

policyTypes:

- Ingress

ingress:

- from:


- podSelector:

matchLabels:

app: frontend

ports:

- protocol: TCP





→ Container Network Interface

Kindnet → Network Policies

Feature	CNI (Container Network Interface)	kube-proxy
Primary Role	Manages pod networking (IP allocation & routing)	Manages Service networking (routing & load balancing)
Scope	Handles Pod-to-Pod communication	Handles Service-to-Pod communication
How It Works	Sets up network interfaces, assigns IPs, configures routes	Uses iptables/ipvs to route traffic to the correct pod
Used For	Pod networking, enforcing network policies (if supported)	Service discovery and load balancing
Examples	Flannel, Calico, Cilium, Weave	Built-in Kubernetes component ( kube-proxy )
Dependency	Required for Kubernetes networking	Optional with advanced CNIs (e.g., Cilium can replace kube-proxy )
Load Balancing	Not responsible for load balancing	Provides basic load balancing for services
Handles Network Policies?	Some CNIs (Calico, Cilium) support them	Does not handle network policies

→ Created NS


→ web app deployment

→ expose it using a service

2 deployments .

allowed

denied



kind: Cluster

apiVersion: kind.x-k8s.io/v1alpha4

networking:

disableDefaultCNI: true # Disables Kindnet

podSubnet: "192.168.0.0/16" # Ensure compatibility with Calico

nodes:

- role: control-plane

- role: worker

kind create cluster --config kind-calico.yaml --name

**\*\*Install Calico\*\***

---

```
kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.26.1/manifests/calico.yaml
```

---

**\*\*We will create a NAMESPACE.\*\***

---

```
kubectl create namespace network-policy-lab
```

---

**\*\*Creating Web Application Deployment\*\***

---

```
kubectl create deployment web-app --image=nginx -n network-policy-lab
```

---

```
kubectl expose deployment web-app --port=80 --target-port=80 -n network-policy-lab
```

---

**\*\*Verify the deployment and Service\*\***

---

```
kubectl get deployments -n network-policy-lab
```

---

```
kubectl get services -n network-policy-lab
```

---

**\*\*Deploy 2 pods using deployment which we will try to connect to web application.\*\***

---

```
kubectl create deployment busybox-allowed --image=busybox -n network-policy-lab -- /bin/sh -c "sleep 3600"
```

---

```
kubectl create deployment busybox-denied --image=busybox -n network-policy-lab -- /bin/sh -c "sleep 3600"
```

---

**\*\*Verify the deployments\*\***

```
kubectl get deployments -n network-policy-lab
```

**\*\*Without Creating any Network Policies lets try to communicate with the web application\*\***

```
kubectl exec -it podname --sh
```

```
wget -qO- http://web-app
```

Able to communicate using both the pods.

**\*\*Now we create a Network Policy\*\***

It controls **\*\*Ingress traffic\*\***, meaning it specifies which pods are **\*\*allowed\*\*** to send traffic to the selected pods.

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: allow-busybox-allowed
```

```
  namespace: network-policy-lab
```

```
spec:
```

podSelector:

matchLabels:

app: web-app

policyTypes:

- Ingress

ingress:

- from:

- podSelector:

matchLabels:

app: busybox-allowed

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: deny-all-ingress

namespace: network-policy-lab

spec:

podSelector:

matchLabels:

app: web-app

policyTypes:

- Ingress