

# Kubernetes state Persistence

starts at 9:05pm

## Agenda

1. Volumes in Kubernetes

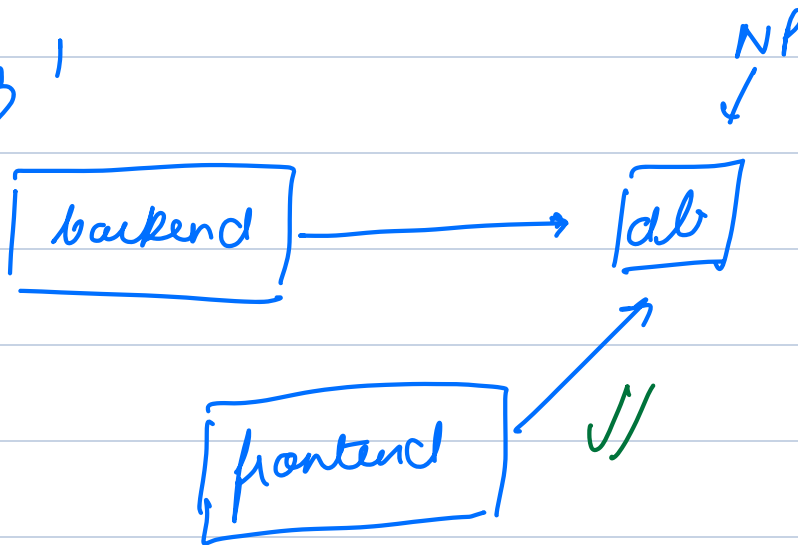
2. Persistent Volumes and Persistent Volume Claims

3. Storage Classes


4. Stateful Sets

1. Storage in Stateful sets

→ Quiz 1



→ Volumes in K8s



container restarts	ES → lost
pod rescheduled	ES → lost
pod is deleted	" "

to persist data → Volumes.


### Types of Volumes

① EmptyDir { }

exists as long as the pod exists.  
pod is deleted → data is lost.

Why emptyDir { }

- ① Persistence across container restarts.
- ② sharing data b/w containers.



③ Performance. → Fast storage.

```
emptyDir: { medium: "Memory" }
```

*emptyDir: 1. }*

Feature	Writing to Container Filesystem	Using emptyDir Volume
Data Persistence	Lost when container restarts	Survives container restarts (but lost when pod is deleted)
Shared Storage	No, isolated per container	Yes, shared across containers in a pod
Performance	Uses container's disk storage	Can be <b>memory-backed</b> for high speed
Pod Deletion	Data is lost	Data is lost

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: myapp
```

```
  labels:
```

```
    app: myapp
```

```
spec:
```

```
  containers:
```

```
    - name: myapp
```

```
      image: alpine:latest
```

```
      command: ['sh', '-c', 'while true; do echo "logging $(date)" >> /opt/logs.txt; sleep 1; done']
```

volumeMounts:

- name: data

mountPath: /opt

- name: logshipper

image: alpine:latest

command: ['sh', '-c', 'tail -F /opt/logs.txt']

volumeMounts:


- name: data

mountPath: /opt

volumes:

- name: data

emptyDir: {}




## ② Configmap and Secrets as Volumes.

key 1 : value 1  
key 2 : value 2.



Pod.




/tmp → mount path.

→ key 1 → file 1

→ key 2 → file 2

>

③ hostpath.



apiVersion: v1

kind: Pod

metadata:

name: hostpath-demo

spec:

containers:

- name: busybox

image: busybox

command: ["/bin/sh", "-c"]

args:

- while true; do

echo "\$(date) - Log entry from pod" >> /var/log/app-logs/demo.log;

sleep 5;

done

volumeMounts:

- mountPath: /var/log/app-logs

name: host-volume

volumes:

- name: host-volume

hostPath:


path: /var/log/app-logs # Host directory to be mounted

type: DirectoryOrCreate # Creates directory if not exists

## Limitations.

- ① Pods can modify the host file.
- ② Pods must be scheduled on the same node.
- ③ Breaks Portability.

## Persistent Volumes.



→ Pods can be rescheduled on other nodes.

A \*\*Persistent Volume (PV)\*\* is a \*\*pre-provisioned storage resource\*\* in the cluster, managed by an administrator.

→ storage capacity

→ access modes

→ reclaim policy.

→ SOGi

→ Access Modes .

## ① Read Write Once (RWO)

- The volume can be mounted as read-write by a single node.

- **Only one pod** can use it at a time

## ② Read Only Many (ROX)

- The volume can be mounted as read-only by many nodes.

- Multiple pods can **read** the volume at the same time


- No write access

## ③ Read Write Many (RWX)

- Multiple pods can read and write simultaneously


- Works **across multiple nodes**

Access Mode	Symbol	Read/Write	Multiple Pods?	Multiple Nodes?
ReadWriteOnce	RWO	Read/Write	✗ No	✗ No
ReadOnlyMany	ROX	Read-Only	✓ Yes	✓ Yes
ReadWriteMany	RWX	Read/Write	✓ Yes	✓ Yes





# Persistent Volume Claim.



## \*\*Binding Process\*\*:

- When a PVC is created, Kubernetes looks for a PV that matches the PVC's request based on the storage class, size, and access modes.
- Once a suitable PV is found, it is bound to the PVC.
- This binding process makes the storage available to the Pod that requested it.

## Reclaim Policies. →

### ① Retain

PVC → deleted

PV → Retained

↓  
Released?

② Delete.

PVC → deleted

PV → deleted.

EBS  
Azure files

deleted.

Reclaim Policy	What It Does?	Common Use Cases
Retain (default)	Keeps the PV and its data, even if the PVC is deleted	Critical data (e.g., databases, logs)
Delete	Deletes the PV along with its storage (cloud disk, EBS, GCE)	Cloud storage (Amazon EBS, GCE Persistent Disks)

→ Demo (Read write Once)

Feature	ReadWriteOnce (RWO)	hostPath
<b>Data Storage Location</b>	Stored in a <b>Persistent Volume (PV)</b> (e.g., EBS, Azure Disk, etc.)	Stored directly on the <b>host node's filesystem</b>
<b>Pod Scheduling</b>	Kubernetes ensures the pod runs on the node where the PV is attached	Pod must be manually scheduled on the same node
<b>Security &amp; Isolation</b>	More secure, managed storage	Less secure, pod can access critical host files
<b>Persistence</b>	Data persists even if the node fails (depending on storage backend)	Data is lost if the node is replaced

apiVersion: v1

kind: PersistentVolume

metadata:

name: pv-cloud

spec:

capacity:

storage: 10Gi

accessModes:

- ReadWriteOnce

persistentVolumeReclaimPolicy: Retain

awsElasticBlockStore:

volumeID: vol-1234567890abcdef

fsType: ext4

Break → 10:40 pm.

→ Demo

**\*\*Create a Persistent Volume with Retain Policy (RWO)\*\***

apiVersion: v1

kind: PersistentVolume

metadata:

name: pv-retain

spec:

capacity:

storage: 5Gi

accessModes:

- ReadWriteOnce

persistentVolumeReclaimPolicy: Retain

storageClassName: standard

hostPath:

path: "/mnt/data-retain" # This path is on your node

If this PV is used by a pod, the pod will see **\*\*the same files stored in `/mnt/data-retain` on the host\*\***.

**\*\*Create a Persistent Volume Claim\*\***

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: pvc-retain

spec:

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 5Gi

**\*\*Create a Pod\*\***

apiVersion: v1

kind: Pod

metadata:

name: pvc-test

spec:

volumes:

- name: storage

persistentVolumeClaim:

claimName: pvc-retain

containers:

- name: test-container

image: busybox

command: ["sleep", "3600"]

volumeMounts:

- mountPath: "/data"

name: storage

→ Limitations of PV

- ① Manual Provisioning.
- ② Lack of Dynamic Allocation.
- ③ Waste of Resources.
- ④ Retention Policy limitation.


50 Gi → PV

10 Gi → PVC

40 Gi → ?  
          

→ Storage Class

100 GB



PVC

70 GB

A **StorageClass** in Kubernetes provides a way to dynamically provision **Persistent Volumes (PVs)** on demand.

PV → ? No.

PVC → PV will be created

PVC 20GB ? → PV 20GB.

Key Components of storage class.

1. **Provisioner** → Defines how Kubernetes interacts with the storage provider.

2. **Parameters** → Defines backend-specific options (like disk type, IOPS, etc.).

3. **Reclaim Policy** → What happens when a PVC is deleted? ('Retain', 'Delete').

- **Retain** → PV stays after PVC deletion. Manual cleanup required.



- **Delete** → PV is deleted automatically when PVC is deleted.

1. **Volume Binding Mode** → Controls when and how PVs are bound (Immediate, WaitForFirstConsumer).

1. When **volumeBindingMode: Immediate** is used, the **Persistent Volume (PV)** is provisioned and bound to the PVC as soon as the PVC is created, even if no pod is using it yet.

Provisioner	Storage Backend
<code>ebs.csi.aws.com</code>	AWS Elastic Block Store (EBS)
<code>pd.csi.storage.gke.io</code>	Google Persistent Disk
<code>disk.csi.azure.com</code>	Azure Disk
<code>nfs.csi.k8s.io</code>	NFS (Network File System)
<code>hostpath</code>	Host machine storage
<code>rancher.io</code>	Kind

→ immediate (Volume Binding mode)  
→ wait for first consumer.





---

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: fast-storage

provisioner: ebs.csi.aws.com # Provisioner for AWS EBS

parameters:

type: gp3 # AWS EBS Volume Type

iops: "3000"

throughput: "125"

reclaimPolicy: Retain # PVs remain even after PVC is deleted

volumeBindingMode: WaitForFirstConsumer # PV is created only when a Pod uses the PVC

allowVolumeExpansion: true # PVCs can request more storage later

---

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: expandable-pvc

spec:

accessModes:

- ReadWriteOnce

resources:

---

requests:

storage: 10Gi

storageClassName: fast-storage

```
kubectl patch pvc expandable-pvc -p '{"spec":{"resources":{"requests":{"storage":"20Gi"}}}}'
```

#### - **File Storage (Default)**

- `volumeMode: Filesystem` (default if not specified).

- The PV is formatted with a filesystem (e.g., ext4, xfs).

- Pods access it via a **file system interface**.

- Examples: **NFS, EFS (AWS), CephFS, HostPath**.

#### - **Block Storage**

- `volumeMode: Block`.

- The PV is **raw block storage** without a filesystem.

- Pods access it **as a raw device** (like `/dev/xvdb`).


- Examples: **AWS EBS, GCE Persistent Disk, OpenStack Cinder**.

Reclaim Policy → Delete.  
volume Binding Mode → Wait for first  
customer.

3 Gi



Pod is deleted



PV → deleted.

PV ? PVC ?

PV will remain

PVC → pending.

StorageClass Demo.

**\*\*Storage Class\*\***

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: my-storage

provisioner: rancher.io/local-path # correct provisioner for Kind

reclaimPolicy: Retain

volumeBindingMode: WaitForFirstConsumer

**\*\*PersistentVolumeClaim\*\***

```
apiVersion: v1

kind: PersistentVolumeClaim

metadata:

  name: my-pvc

spec:

  accessModes:

    - ReadWriteOnce

  resources:

    requests:

      storage: 2Gi

  storageClassName: my-storage
```

**\*\*Pod will use the pvc which will create the PV\*\***

```
apiVersion: v1

kind: Pod

metadata:

  name: my-pod

spec:

  volumes:

    - name: storage-volume

      persistentVolumeClaim:

        claimName: my-pvc
```

containers:

- name: app

image: busybox

command: [ "sleep", "3600" ]

volumeMounts:

- mountPath: /data

name: storage-volume