

# Kubernetes Security

starts at 9:05pm

## Agenda

1. Stateful sets

2. Authentication

1. Types

3. Authorisation

1. Types


4. Demo Authentication Authorisation

5. Security Contexts

6. Admission Controller

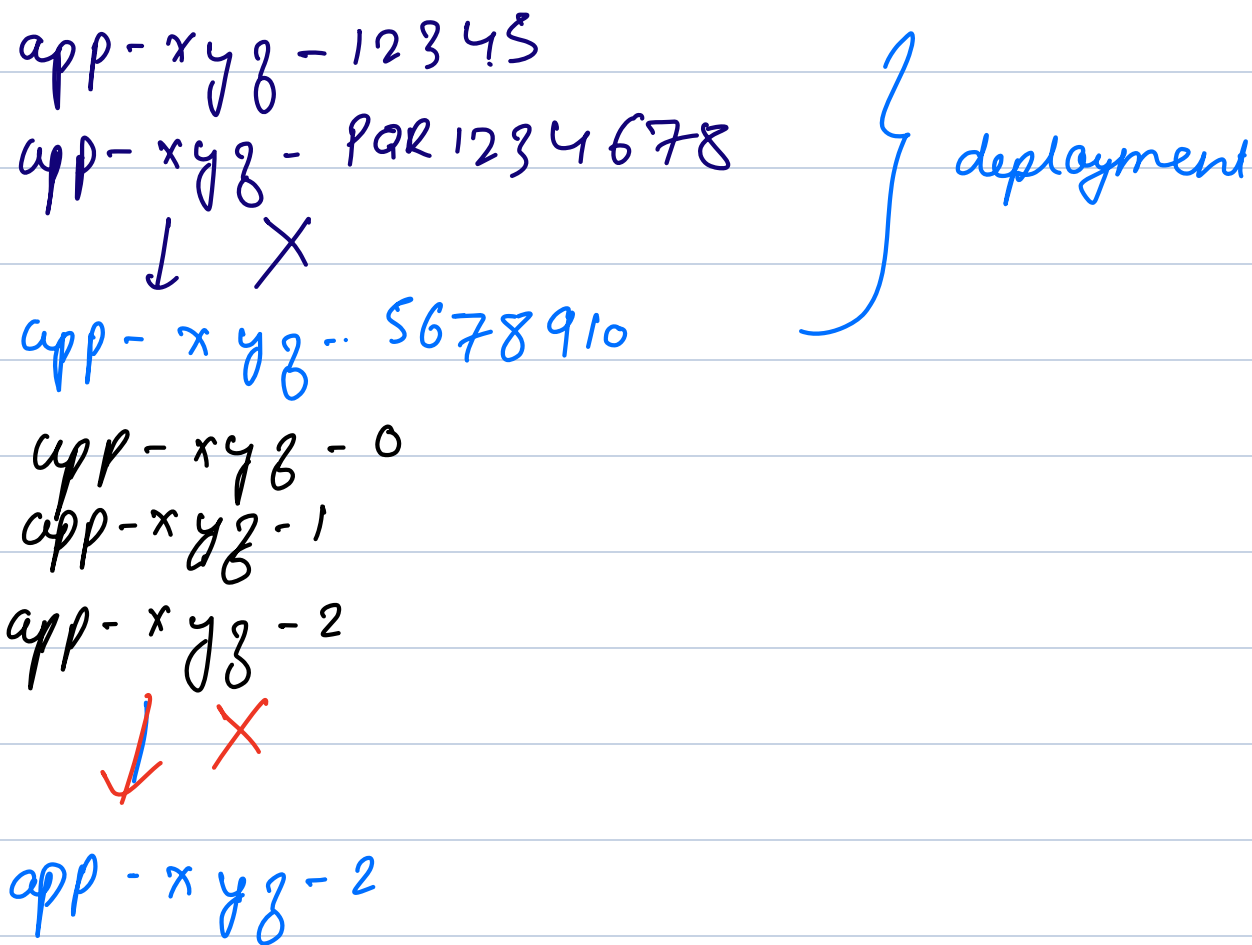
## → Stateful sets


A **stateful application** is an application that **maintains its state** across different sessions, meaning it remembers data and context from one operation to the next.



A **StatefulSet** is a Kubernetes workload API object used to **manage stateful applications**.

Feature	Deployment	StatefulSet
Use Case	Stateless applications	Stateful applications (DBs, etc.)
Pod Identity	No stable identity, random names	Stable, unique identities
Storage	Ephemeral or shared storage	Persistent storage with PVCs
Scaling	Pods scaled randomly	Pods scaled sequentially
Service Type	Normal Service (for load balancing)	Headless Service (for unique DNS)





volumeClaimTemplates:

- metadata:

name: mysql-data

spec:

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 5Gi

storageClassName: standard

## Demo

### ① Headless Service -

my-db-0.db-service.my-namespace1.svc.cluster.local

apiVersion: v1

kind: Service

metadata:

name: db-service

spec:

clusterIP: None # This makes the service headless (no IP assignment)

selector:

app: my-db

ports:

- port: 5432

targetPort: 5432

*Stateful set (different pv, pvc's per each pod)*

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: my-db

spec:

serviceName: db-service # Link to the headless service

replicas: 3 # Number of replicas (PostgreSQL pods)

selector:

matchLabels:

app: my-db

template:

metadata:

labels:

app: my-db

spec:

containers:

- name: database

image: postgres

env:

- name: POSTGRES\_PASSWORD

value: "mysecretpassword" # Set the password for PostgreSQL

ports:

- containerPort: 5432

volumeMounts:

- name: db-storage

mountPath: /var/lib/postgresql/data # Path where the DB will store data

volumeClaimTemplates:

- metadata:

name: db-storage

spec:

accessModes: [ "ReadWriteOnce" ]

resources:

requests:

storage: 5Gi

## Kubernetes Security.

Authentication

Authorization


Admission Controller.

### Authentication.

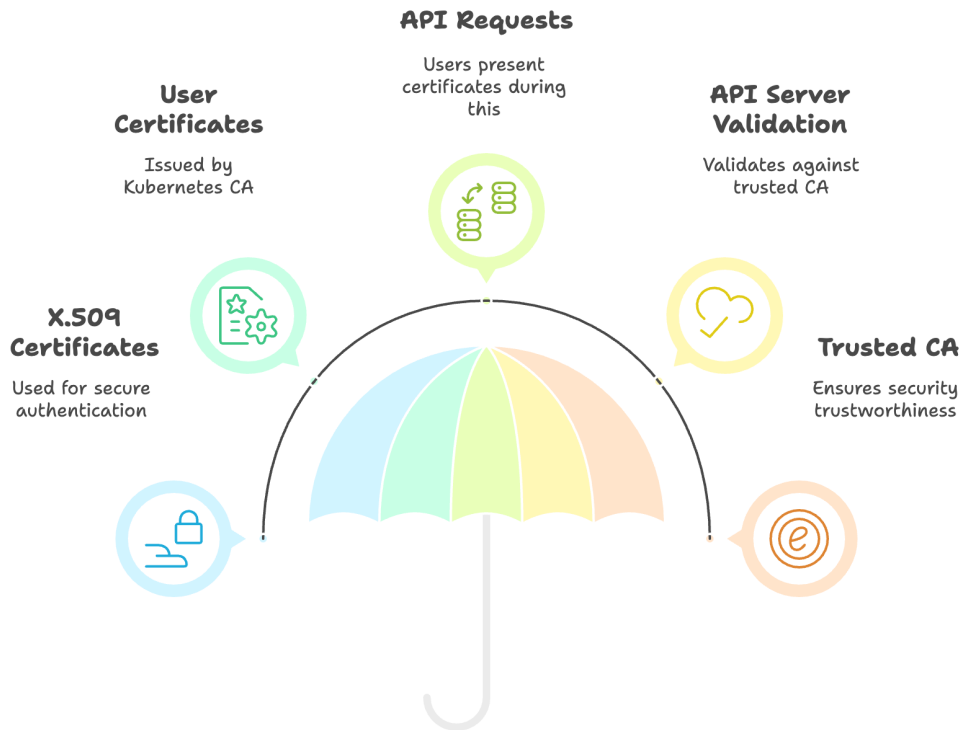
\*\*Authentication is the process of verifying the identity of users and systems that interact with the Kubernetes cluster.\*\*

### Different authentication Methods.

#### ① Client Certificates.




## Kubernetes Authentication Process



② Bearer Token Authentication.

① Service Account Token.

→ Pod.



`/var/run/secrets/kubernetes.io/serviceaccount/token`

This token allows the Pod to authenticate with the API Server.

## ② Static Token file.

user  $\rightarrow$  SA  $\rightarrow$  grant API access.

/etc/kubernetes/static\_tokens.csv

Token never expires



Security Risk.

```
echo "my-static-token,user1,user1,system:masters" > /etc/kubernetes/static_tokens.csv
```

## ③ OpenID Connect Authentication.

It enables Kubernetes to authenticate users using external identity providers.

eg:- Okta?

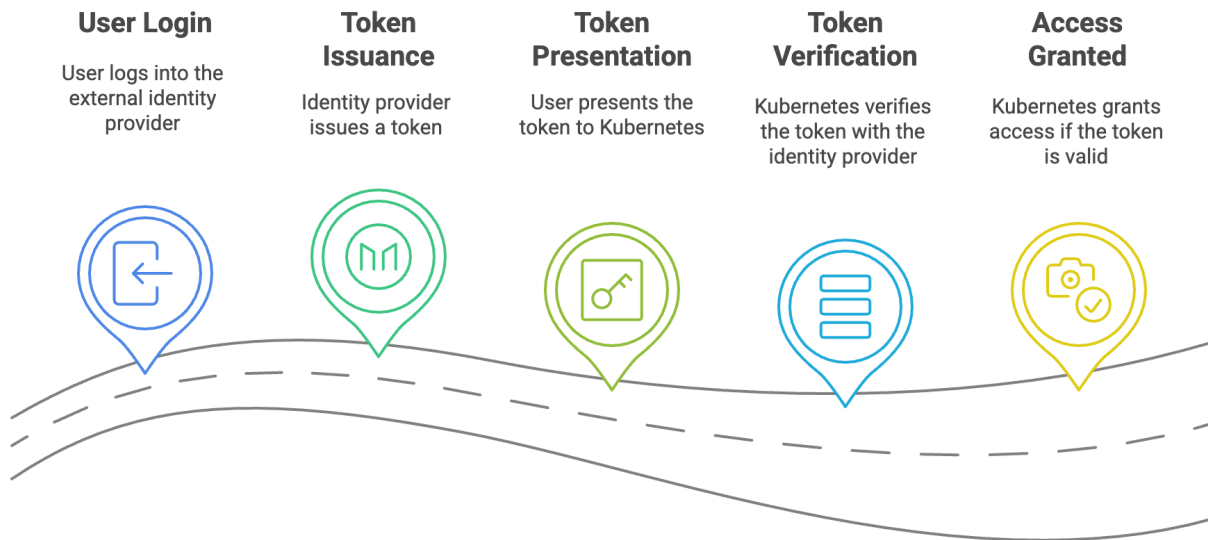
Google

Azure.


IPd.



## Kubernetes Token-Based Authentication Process

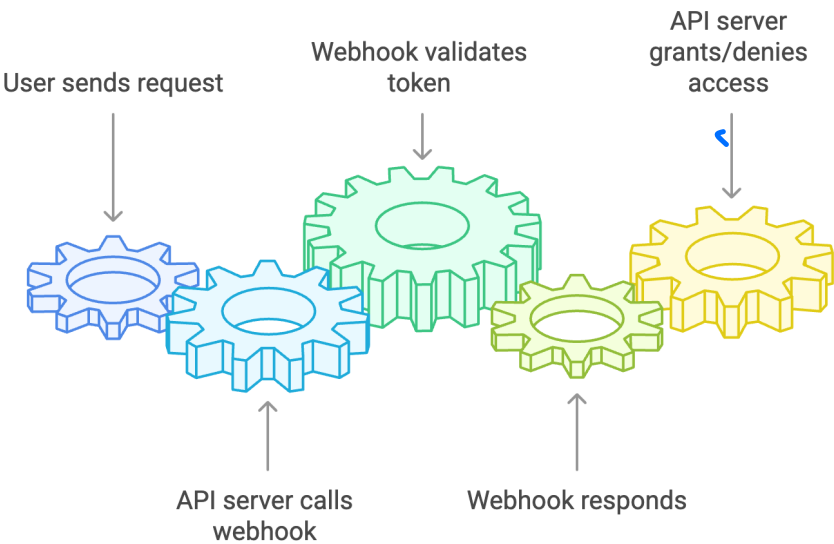


## ④ Webhook Token Authentication.



Webhook authentication allows Kubernetes to **\*\*delegate authentication\*\*** to an external service.

Kubernetes API Authentication Process



Authentication Method	How It Works	Use Case	Pros	Cons
Client Certificates	Uses X.509 certificates to authenticate users and systems.	Secure authentication for admins and system processes.	Strong security, cryptographic validation.	Requires <b>certificate management</b> .
Service Account Tokens	Kubernetes generates tokens for Pods to access the API.	Used by applications running inside the cluster.	Automatic, easy to use.	Token rotation and security risks if leaked.
Static Token File	Stores predefined tokens in a file on the API server.	Used for <b>testing or simple setups</b> .	Easy setup.	<b>Insecure</b> , manual management.
OIDC Authentication	Uses an external identity provider (Google, Azure, Okta).	Enterprise authentication with <b>corporate credentials</b> .	Centralized identity management, easy integration.	Requires an external provider and setup.
Webhook Token Authentication	API server calls an external webhook to validate tokens.	Custom authentication using a separate service.	<b>Flexible</b> , integrates with external systems.	Requires <b>maintaining a separate authentication service</b> .

Break → 10:20pm.

Authorisation.

→ authenticated → what all can  
you do  
inside the cluster?

**\*\*Authorization is used for controlling access to resources\*\***

Different Authorisation Modes.

① Node.

used by kubelet.

This mode authorizes kubelets to read and write to their node objects.

② ABAC.

Attribute Based Access Control.

{

"apiVersion": "abac.authorization.kubernetes.io/v1beta1",

"kind": "Policy",

```
"spec": {
```

```
"user": "alice",
```

```
"namespace": "default",
```

```
"resource": "pods",
```

```
"readonly": true
```

```
}
```

```
}
```

```
- --authorization-mode=Node,RBAC,ABAC
```

```
- --authorization-policy-file=/etc/kubernetes/abac-policy.json
```

After adding abac-policy.json



Restart the API server pod.

**\*\*Why is ABAC Less Preferred?\*\***

- Policies are stored in a **\*\*static file\*\***, requiring API server restarts for updates.

- **\*\*Not scalable\*\*** for large clusters.

- **\*\*RBAC is preferred\*\*** because it is **\*\*dynamic and easier to manage\*\***.

### ③ Webhook Authorization.

A **custom external service** makes authorization decisions. When a request is received:

- 1 . Kubernetes sends it to an external **webhook server**.
- 2 . The server **verifies the request** based on custom logic.
- 3 . The server **returns allow/deny decisions**.

```
- --authentication-token-webhook-config-file=/etc/kubernetes/webhook-config.yaml
```

```
- --authentication-token-webhook-cache-ttl=5m
```

```
apiVersion: v1
```

```
kind: Config
```

```
clusters:
```

```
- name: webhook-auth
```

```
  cluster:
```

```
    server: https://auth.example.com/validate
```

```
users:
```

```
- name: webhook-auth
```

```
contexts:
```

```
- context:
```

```
  cluster: webhook-auth
```

```
  user: webhook-auth
```

```
  name: webhook-auth
```

→ RBAC

Roles → Permissions  
Rolebindings → Assigned to users  
that make use of this  
job.

Roles. → within a namespace  
ClusterRoles. → across cluster.

kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

namespace: default

name: pod-reader

rules:

- apiGroups: [""]

resources: ["pods"]

verbs: ["get", "watch", "list"]

API Group	Example Resources
"" (empty, core API)	pods, services, configmaps, secrets, nodes, namespaces
apps	deployments, daemonsets, statefulsets, replicaset
batch	jobs, cronjobs
rbac.authorization.k8s.io	roles, rolebindings, clusterroles, clusterrolebindings
networking.k8s.io	networkpolicies, ingresses
storage.k8s.io	storageclasses, volumeattachments

Verb	Description
get	Read a specific resource (e.g., <code>kubectl get pod my-pod</code> )
list	List all resources of a type (e.g., <code>kubectl get pods</code> )
watch	Monitor changes to resources in real-time
create	Create a new resource
update	Modify an existing resource
patch	Partially update an existing resource
delete	Remove a resource
deletecollection	Delete multiple resources at once

*Role Bindings* Associates role with a user  
*ClusterRole Bindings.* or group.

apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

name: read-pods

namespace: default

subjects:

- kind: User

Jane → Pod-reader

name: "jane"

apiGroup: rbac.authorization.k8s.io

roleRef:

kind: Role

name: pod-reader

apiGroup: rbac.authorization.k8s.io

subjects:

- kind: User

name: "jane"

apiGroup: rbac.authorization.k8s.io

- kind: Group

name: "dev-team"

apiGroup: rbac.authorization.k8s.io

- kind: ServiceAccount

name: "app-sa"

namespace: default

roleRef:

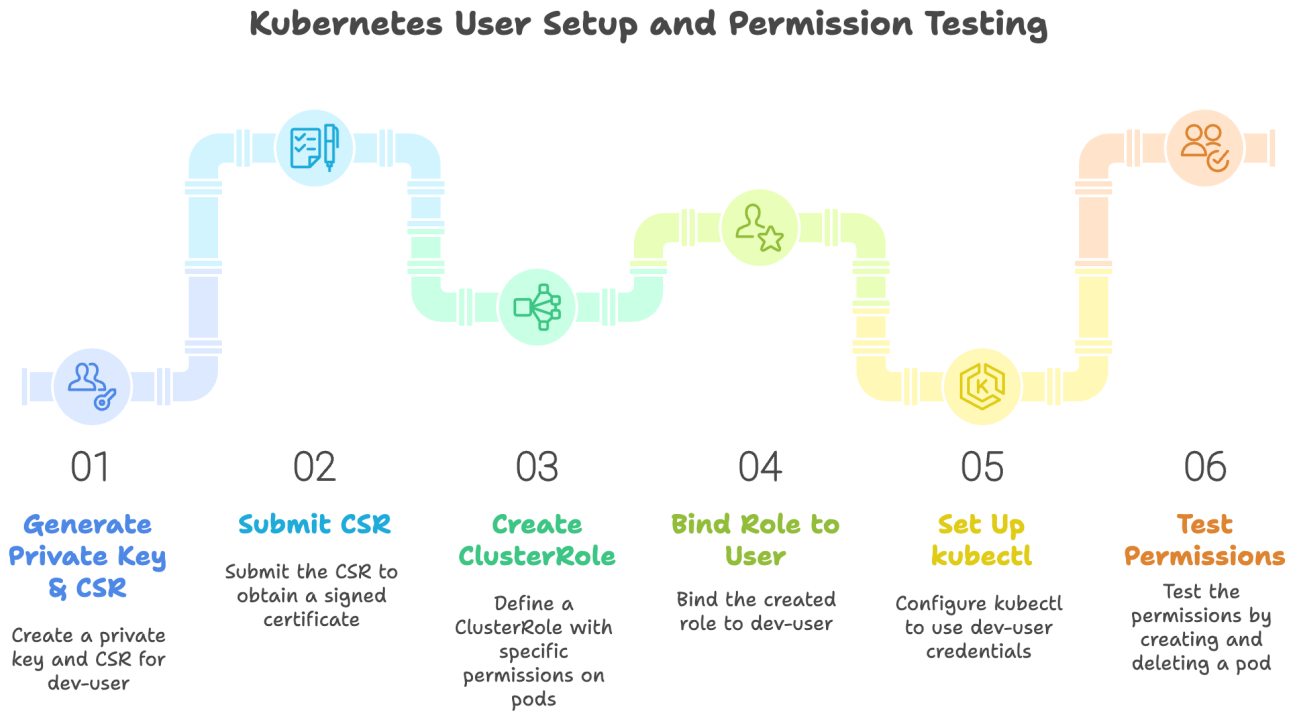
kind: Role



name: pod-reader

apiGroup: rbac.authorization.k8s.io

→ Authentication and Authorisation Demo.



dev user → Authentication.

→ Key  
↓  
→ CSR  $\xrightarrow[\text{CA}]{\text{Submit to}}$  Certificate.

kube config.



certificate & key (dev-user)



cluster?

Context.

cluster + user.

→ Security Contexts

Security contexts are settings applied to Pods and containers to define privilege and access control settings.

→ Pod level

→ Container level -

Pod level

apiVersion: v1

kind: Pod

metadata:

name: secure-pod

spec:

securityContext: # Security settings apply to all containers in the Pod

runAsUser: 1000

runAsGroup: 1000

fsGroup: 1000

containers:

- name: app-container

image: nginx

*Container level*

apiVersion: v1

kind: Pod

metadata:

name: secure-container

spec:

containers:

- name: app-container

image: nginx

securityContext: # Only applies to this container

runAsUser: 2000

readOnlyRootFilesystem: true

*Key Security Context Fields*

*① runAsUser & runAsGroup.*

securityContext:

runAsUser: 1000 # Runs as non-root user with UID 1000

runAsGroup: 3000 # Uses GID 3000

② Privileged.

securityContext:

privileged: true # Grants full access to the host

③ Allow Privilege Escalation.

Blocking Privilege Escalation

securityContext:

allowPrivilegeEscalation: false

④ Read Only Filesystem.

securityContext:

readOnlyRootFilesystem: true

⑤ Linux capabilities

---

securityContext:

capabilities:

add: ["NET\_ADMIN"] # Allows changing network settings

drop: ["ALL"] # Drops all unnecessary capabilities

apiVersion: v1

kind: Pod

metadata:

name: secure-pod

spec:

securityContext: # Applies to all containers in the pod

runAsUser: 1000 # Runs as non-root user

runAsGroup: 3000 # Group ID for security

fsGroup: 2000 # Ensures correct file permissions

containers:

- name: secure-container

image: busybox

securityContext:

allowPrivilegeEscalation: false # Prevent privilege escalation

readOnlyRootFilesystem: true # Prevents modifying root FS

command: ["sleep", "3600"]

---

```
kubectl exec -it secure-pod -- id
```

```
kubectl exec -it secure-pod -- sh -c "su root"
```

```
kubectl exec -it secure-pod -- touch /testfile
```