

Kubernetes API Versioning, K8s Extension, Certification Tips & Helm Introduction.

starts at 9:05 pm

AGENDA

1. Admission Controller

1. Types

2. API Version

3. CRD / CR

4. CKAD Exam tips

1. Tips

2. Killer.sh Overview

3. Sample Questions

spec:

securityContext:

readOnlyRootFilesystem: true

containers:

- name: app

image: nginx

Admission Controller

→ validate

→ modify or

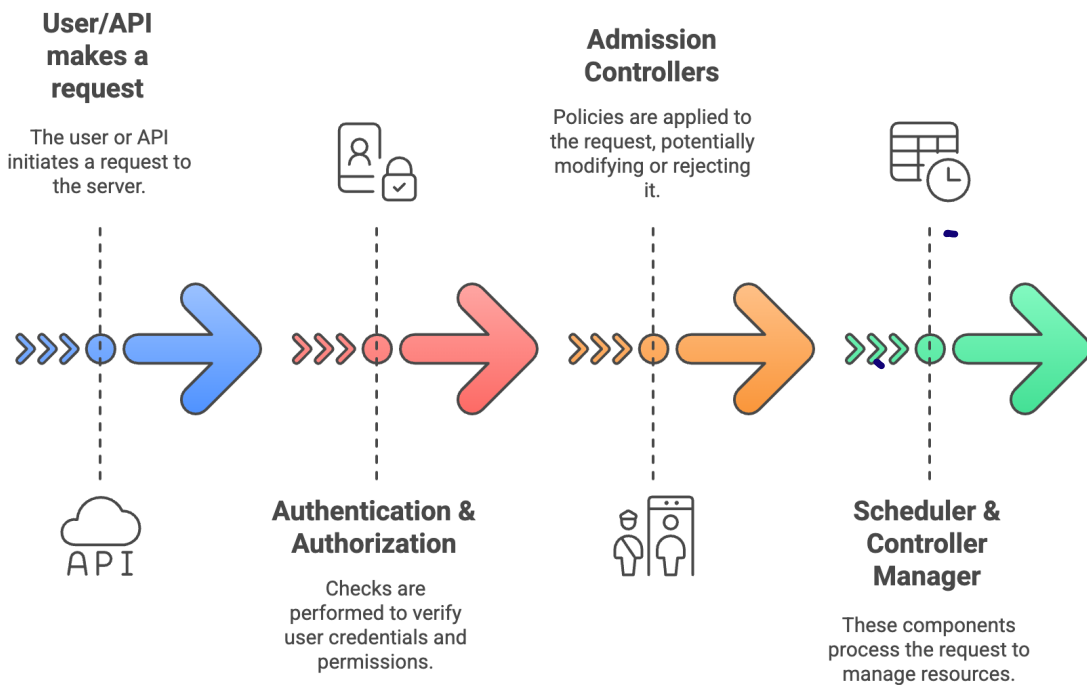
→ reject API request.

→ before getting applied to the cluster.

****They intercept requests to the Kubernetes API server before an object is created,****

****but after the request is authenticated and authorized.****

API Request Handling in Kubernetes



Types of Admission Controllers.

- ① Mutating
- ② Validating


↓
x can't modify object.


Resource Quota ?

↓
validating AC.


→ Mutating AC

Admission Controller	Description
MutatingAdmissionWebhook	Calls an external webhook to modify API requests before storing them.
AlwaysPullImages	Forces all pods to pull images from the registry instead of using cached images.
DefaultStorageClass	Automatically assigns a default storage class to Persistent Volume Claims (PVCs) if none is specified.
LimitRanger	Adds default resource requests and limits if they are missing in a pod/container specification.





No requests or limits



Validating AC

Admission Controller	Description
ValidatingAdmissionWebhook	Calls an external webhook to validate API requests before accepting them.
ResourceQuota	Enforces namespace-level quotas (CPU, memory, object count).
NodeRestriction	Restricts kubelet from modifying critical node and pod objects.

→ Dual Purpose. (Both mutating and validating)

Limit Range →

Adds default resource requests/limits if missing (mutates) and rejects pods that exceed limits (validates).

→ Demo

```
docker exec -it admission-demo-control-plane bash -c \
```

```
"sed -i 's/\(--enable-admission-plugins=[^ ]*\)/\1,LimitRanger,AlwaysPullImages/' /etc/kubernetes/manifests/kube-apiserver.yaml"
```

↘ restart Api-server after this
↓
delete the pod.

```
apiVersion: v1
```

```
kind: LimitRange
```

```
metadata:
```

```
  name: pod-resource-limits
```

```
  namespace: limitranger-demo
```

```
spec:
```

```
  limits:
```

```
    - type: Container
```

```
      default:
```

```
        cpu: "500m"
```

```
        memory: "256Mi"
```

```
      defaultRequest:
```

```
        cpu: "250m"
```

```
        memory: "128Mi"
```

→ API versions

v1 v2 v3 . . .

deprecated (still available → not suggested)
removed. (not compatible)

→ apiVersion: group / version.

API groups.

API Group	Description	Example Resources
"" (core API, empty string)	Core API with fundamental resources	Pods, nodes, services, secrets, configmaps
apps	Manages workloads and deployments	deployments, statefulsets, daemonsets
batch	Handles batch job workloads	jobs, cronjobs
networking.k8s.io	Controls networking-related resources	networkpolicies, ingress, ingressclasses
storage.k8s.io	Manages storage provisioning	storageclasses, volumeattachments
certificates.k8s.io	Handles certificate signing requests (CSRs)	certificatesigningrequests (CSR)
admissionregistration.k8s.io	Controls admission webhooks	mutatingwebhookconfigurations, validatingwebhookconfigurations

Alpha

Beta

Stable.

| **Alpha** | Experimental, may change anytime, not recommended for production. |

| ----- | ----- |

| **Beta** | More stable, but still subject to change, recommended for testing. |

| **Stable (GA)** | Fully supported and stable, safe for production use. |

Evolution of Deployments

1.8 → apps/v1 beta1

1.9 → apps/v1 beta2

1.10 → apps/v1

1.15

Removed?

networking.k8s.io/v1beta1

V. 1.19 → Kubernetes.

(deprecated)

V. 1.22 Kubernetes.
(Removed.)

V 1.24

apiVersion: networking.k8s.io/v1beta1

kind: Ingress

metadata:

name: my-ingress

namespace: default

spec:

rules:

- host: my-app.example.com

http:

paths:

- path: /

backend:

serviceName: my-service

servicePort: 80

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: my-ingress

namespace: default

spec:

rules:

- host: my-app.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: my-service

port:

number: 80

1.24

kubectl convert

break → 10:15 pm

Custom Resource Definition & Custom Resources

Science → Books

Maths → "


History. → "

" " → Pod

apps → Deployment

batch → job, cronjob.

Science Fiction.?



A **Custom Resource Definition (CRD)** is a way to extend Kubernetes with your **own custom resources**, just like built-in ones (Pods, Deployments, Services, etc.).

(Template) → CRD.

It **defines the structure** of a new resource type that Kubernetes should recognize.

v 1.10 → Ingress didn't exist

custom nginx, haproxy.


- **Kubernetes v1.1 (2015):** Ingress was **introduced as a built-in resource** under `networking.k8s.io/v1beta1`.

- **Kubernetes v1.19 (2020):** Ingress **graduated to GA** as `networking.k8s.io/v1`.

CRD → fruits

CR → apple.

Horizontal Pod Autoscaler. Autoscaling.



CRD

apiVersion: apiextensions.k8s.io/v1

kind: CustomResourceDefinition

metadata:

name: fruits.mygroup.example.com

spec:

group: mygroup.example.com

scope: Namespaced

names:

plural: fruits

singular: fruit

kind: Fruit

shortNames:

- fr

versions:

- name: v1

served: true

storage: true

schema:

openAPIV3Schema:

type: object

properties:

spec:

type: object

properties:

color:

type: string

taste:

type: string

`metadata.name` → The **fully qualified name** of the CRD.

- Follows the format: **<plural>.<group>** → `fruits.mygroup.example.com`.

- This is **how Kubernetes uniquely identifies this CRD**.

`group: mygroup.example.com` → Defines the **API group** where this resource will belong.

`scope: Namespaced` → Means that **each "Fruit" resource will exist within a specific namespace**.

- If set to `Cluster`, it would be **cluster-wide** instead.

- `plural: fruits` → The plural name when interacting with multiple resources (`kubectl get fruits`).

- `singular: fruit` → The singular form (`kubectl get fruit apple`).

- `kind: Fruit` → The **Kubernetes Kind** for this resource (`kubectl get Fruit`).

- `shortNames: ["fr"]` → A shorthand for easier use (`kubectl get fr`).

- `name: v1` → Specifies that this **version of the CRD is v1**.

- `served: true` → Kubernetes **accepts requests** for this version.

- `storage: true` → This version is the **primary version stored in etcd**.

- `schema: openAPIV3Schema` → Defines the **structure of the resource using OpenAPI v3 schema**. (Some standard)

- `type: object` → The CRD is an **object** with multiple fields.

- `properties: spec` → The `spec` field contains the actual **custom fields** of the resource.

- `properties:`

- `color: type: string` → **Defines a "color" field**, which must be a string (e.g., `"red"`).

- `taste: type: string` → **Defines a "taste" field**, which must also be a string (e.g., `"sweet"`).

apiVersion: mygroup.example.com/v1

kind: Fruit

metadata:

name: apple

spec:

color: red

taste: sweet

} Custom Resource.

Tips for CKAD exam

****Create a Pod****

```
kubectl run nginx-pod --image=nginx
```

****Create a Deployment****

```
kubectl create deployment nginx-deployment --image=nginx
```

****Exposing a Deployment****

Creates a Service

```
kubectl expose deployment nginx-deployment --port=80 --target-port=80
```

****Scaling a Deployment****

```
kubectl scale deployment nginx-deployment --replicas=3
```

****Updating a Resource****

Can be Asked.

```
kubectl set image deployment/nginx-deployment nginx=nginx:1.16.1
```

****Deleting a Resource****

```
kubectl delete pod nginx-pod
```

****Using Aliases for Speed****

```
alias k='kubectl'
```

alias kgp='kubectl get pods'

alias kgn='kubectl get nodes'

alias kaf='kubectl apply -f'

****Using Dry-Run to Generate YAML****

kubectl create deployment nginx-deployment --image=nginx --dry-run=client -o yaml > deployment.yaml

****Displays additional information.****

kubectl get pods -o wide

kubectl get pods -o json

kubectl get pods -o yaml

****Filtering Output with `--field-selector`**:**

kubectl get pods --field-selector status.phase=Running

****Using `jq` for JSON Processing**:**

- `jq` is a powerful tool for processing JSON output. You can use it to filter and format JSON data from `kubectl`.

kubectl get pods -o json | jq '.items[] | {name: .metadata.name, namespace: .metadata.namespace,
status: .status.phase}'



Custom Columns:

- The `-o custom-columns` option allows you to specify which columns to display.

```
kubectl get pods -o custom-columns=POD_NAME:.metadata.name,STATUS:.status.phase
```

Using Label Selectors:

```
kubectl get pods -l app=my-app
```

Sorting and Watching Resources:

```
kubectl get pods --sort-by=.metadata.name
```

```
kubectl get pods --watch
```

Using Grep for Quick Searches

```
kubectl get pods -o wide | grep Running
```

What is kill.sh

→ (KAD Question - 1)

In namespace `mct`, create a `PersistentVolume` named `mct-pv`. It should have:

- `Capacity`: 100Mi

- `Access mode`: ReadWriteOnce

- `HostPath`: `/tmp/data`

- `No Storage Class Name` defined

In the same namespace `mct`, create a `PersistentVolumeClaim` named `mct-pvc`. It should:

- `Request Storage`: 100Mi
- `Access Mode`: `ReadWriteOnce`
- `Note`: The PVC should bind to the PV correctly.

In the same namespace `mct`, create a `Pod` named `nginx-mct` using:

- `Image`: `nginx`
- `Container Name`: `mct-container`
- `Mount Path`: `/tmp/projectdata``

`apiVersion: v1`

`kind: PersistentVolume`

`metadata:`

`name: mct-pv`

`namespace: mct`

`spec:`

`capacity:`

`storage: 100Mi`

`accessModes:`

`- ReadWriteOnce`

`hostPath:`

`path: "/tmp/data"`

```
apiVersion: v1

kind: PersistentVolumeClaim

metadata:

  name: mct-pvc

  namespace: mct

spec:

  accessModes:

    - ReadWriteOnce

  resources:

    requests:

      storage: 100Mi
```

```
apiVersion: v1

kind: Pod

metadata:

  name: nginx-mct

  namespace: mct

spec:

  containers:

    - name: mct-container

      image: nginx

      volumeMounts:
```

- mountPath: "/tmp/projectdata"

name: mct-storage

volumes:

- name: mct-storage

persistentVolumeClaim:

claimName: mct-pvc

CKAD question - 2

Developers occasionally need to submit pods that run periodically.

Task

Follow the steps below to create a pod that will start at a predetermined time and which runs to completion only once each time it is started:

Create a YAML formatted Kubernetes manifest that runs the following shell command:

date

in a single **busybox** container. The command should run **every minute** and must **complete within 22 seconds** or be terminated by Kubernetes. The **CronJob name and container name** should both be **hello**.

Create the resource in the above manifest and verify that the job executes successfully at least once.

→ activeDeadlineSeconds → 22

apiVersion: batch/v1

kind: CronJob

metadata:

name: hello

spec:

schedule: "* * * * *" # Runs every minute

jobTemplate:

spec:

activeDeadlineSeconds: 22 # Terminate if not completed within 22s

template:

spec:

containers:

- name: hello

image: busybox

command: ["date"]

restartPolicy: Never

A Kubernetes administrator wants to ensure that all newly created pods specify resource limits before they are scheduled. Which admission controller helps enforce this policy?

A. ResourceQuota

B. **LimitRanger**

C. MutatingAdmissionWebhook

1. Finalisers

2. Topology Constraints

3. DEMO of OIDC

4. Cordon Uncordon

1. Drain Node

5. HPA VPA

6. Istio

7. Linkerd

8. Kyverno

9. Gatekeeper

10. Endpoint Slices

upgrading cluster
Bkp etcd.

Prometheus Rules.